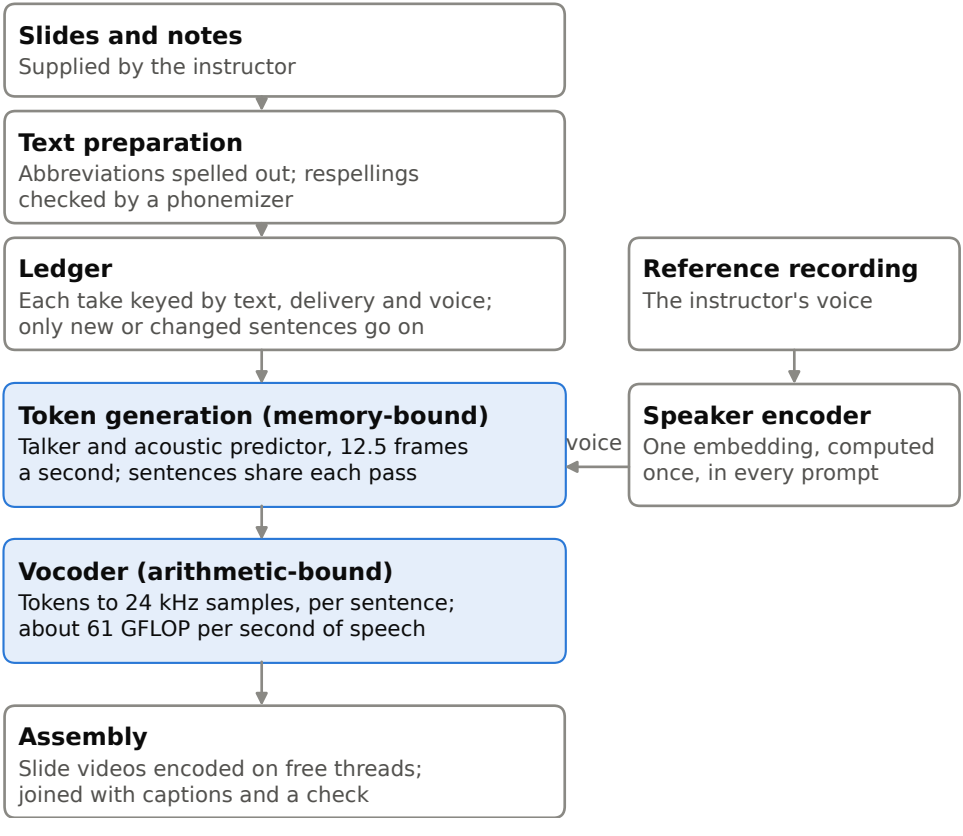
Slides and notes
Supplied by the instructor
Text preparation
Abbreviations spelled out; respellings checked by a phonemizer
Ledger
Each take keyed by text, delivery and voice; only new or changed sentences go on
Reference recording
The instructor's voice
Token generation (memory-bound)
Talker and acoustic predictor, 12.5 frames a second; sentences share each pass
voice
Speaker encoder
One embedding, computed once, in every prompt
Vocoder (arithmetic-bound)
Tokens to 24 kHz samples, per sentence; about 61 GFLOP per second of speech
Assembly
Slide videos encoded on free threads; joined with captions and a check