

A TOOL FOR MAKING LECTURES

Your slides, narrated in your own voice.

Record yourself once. After that, any deck comes back as a finished video — narration, captions, chapters. The voice is cloned on your own laptop, or hosted, as you prefer.

WHAT IT DOES

Three things, one recording

Lectures

A PDF or PowerPoint deck becomes a narrated video with captions and chapter marks. Revise the deck later and only what changed is re-recorded.

Documents

Any text — notes, a paper, a handout — read aloud in your voice, with a caption file beside it.

Conversations

A topic or a paper becomes a two-voice podcast: you and a second speaker, with a greeting and a sign-off.

Three voices: a free draft voice, your own cloned here and sending nothing away, and your own cloned at Google — faster, and closer to you. Whichever speaks it, you read and edit the drafted narration first.

BEFORE YOU START

What you need

Half an hour of setup, most of it unattended.

NEED	WHY	WHERE
Windows 10 or 11	the launchers are .bat files	—
Miniconda	RealMe installs into a conda environment	anaconda.com/download/success
A Gemini API key	drafts the narration, and speaks it	aistudio.google.com
~6 GB free disk	speech engine, model weights, your renders	—
A phone, a quiet room	30-60 seconds of your speech	—

No C++ compiler. No GPU. The local speech engine ships compiled and runs on the CPU; the hosted one needs no engine at all. ffmpeg, the draft voices and the model weights are all in the package.

THE CHOICE THAT MATTERS

Three voices, one recording

One recording clones into either of the two that sound like you.

VOICE	SOUNDS LIKE	COSTS	WHAT LEAVES THE MACHINE
gemini-tts — the default	you	about \$1 an hour	your clip, and every sentence
qwen3cpp	you	hours of CPU	nothing
piper	not you — a draft	nothing	nothing

Draft on piper, record on one of the other two. The hosted voice wants a paid Google project: on the free tier their terms say human reviewers may read what you send, and your voice is personal information.

SETUP

Unzip, install, key, go

Everything is in the package — the speech engine and its weights, three draft voices (Amy, Ryan, Cori), ffmpeg. No compiler, no GPU, no model downloads.

1

Unzip

Put the RealMe folder somewhere permanent. The path does not matter.

2

Install

Double-click 1_Install.bat — the only launcher you need. A few minutes.

3

Key, and check

Set the Gemini key. setup then installs any Python package that is missing.

4

Go

realme studio
Then record your voice, once.

```
cd /d C:\path\to\RealMe
realme key GEMINI_API_KEY your-key-here
realme setup
realme studio
```

No engine step and no update step — the engine travels in the package. Step 2 needs a network connection once, for the Python libraries; everything else is already on disk.

THE ONE-TIME STEP

Record your voice

- Thirty to sixty seconds

Longer does not help. The engine builds a fixed description of your voice either way.

- Read something you have the text of

Paste that exact text in as the transcript — your speaking rate is measured from it.

- One sitting, one room, one microphone

A phone at a steady distance is fine. On iPhone, Voice Memos set to Lossless.

- Read normally

Not slowly, not in a performance voice. The clone reproduces what you give it.

The page tells you whether the recording is worth cloning — length, noise floor, signal-to-noise — before you commit to it.

Save profile

Your voice

Record 30–60 seconds of normal speech — read something you would actually say in a lecture, not a word list. Quiet room, consistent distance from the mic. This gets checked, polished once, and then every lecture inherits the result.

Reference recording

Choose File **bad.wav**

re-record — too much room noise to clone cleanly

Length	14.37s
Peak	-18 dBFS
Noise floor	-35.9 dB
Signal to noise	-2.4 dB

- There is a lot of room in this. Turn off fans and air conditioning, close the window, move away from the computer, and get closer to the mic. Soft furnishings help more than anything you can buy.

Listen and pick one. Conservative is usually right — a clone trained on an over-processed reference sounds processed in every lecture you make, and that cannot be undone later.

► off

 No processing

► clean

 Clean up only — remove noise and rumble, level it

► natural

 Natural polish — clean up, soften sibilance, gentle warmth

► warm

 Warm broadcast — fuller and closer, for a quiet room

Use this voice

THE WORKFLOW

Cheap first, expensive last

Each step costs more than the one above it, and exists so a mistake is caught before the expensive one.

1

Draft

Drop in the deck, describe the course and how you teach. One model call, no audio. Seconds.

2

Edit

Read it. Fix it. This is the step that decides whether the video is any good — and it is free.

3

Hear one line

One button per slide. Checks a pronunciation in seconds, not hours.

4

Render

Leave it running. If it stops, start it again — finished segments are reused.

Two boxes change the result more than any model setting: **who these students are**, and **how you talk**. "Second-year MPH, week 4 of 14. They know regression but not causal diagrams."

IN THE APP

The Lecture tab

The screenshot shows the 'Lecture' tab in the RealMe Studio application. At the top, there are three status indicators: 'MEDIA TOOLS READY' (green), 'NO GEMINI KEY' (yellow), and 'PLACEHOLDER VOICE ONLY' (orange). Below these, the 'Lecture' tab is selected, with 'Podcast & Debate' and 'Twin Setup' also visible. The main content area is titled '1 • Draft the narration' and includes a descriptive paragraph about the step. It features several input fields: 'Slide deck (PDF, PPTX or ODP)' with a 'Choose File' button and 'No file chosen' text; 'Course context' with a text box containing 'PUBH 6002, second-year MPH. They know DAGs but not IPTW. Week 4 of 14.'; and 'Your teaching voice' with a text box containing 'Dry and precise. Worked examples over abstraction. I say 'notice that' a lot and I never read a bullet aloud.' At the bottom, there are two dropdown menus: 'Script writer' set to 'Gemini' and 'Model' set to 'gemini-3-flash (free tier)'. A blue 'Draft narration' button is positioned below these dropdowns. To the right of the main form is an 'Activity' panel showing a single entry: 'Idle.'

The deck goes in at the top. The two context boxes sit beside it. **Draft narration** writes a first version for every slide.

If the deck has speaker notes, the narration is drafted from them. If it has none, it is drafted from the visible slide text — and it reads exactly like that.

Long renders are background jobs. The log survives closing the app.

Editing what will be said

Everything here is free. Every word below it costs minutes of compute, so fix things here first.


- Say the numbers you mean
"15.3 percent" is safer than
"15.3%" in a sentence that also
uses the symbol.
- Acronyms are spelled out
AUC becomes "A U C". A term that
is really a word — NHANES — gets
a lexicon entry.
- Equations are converted
from LaTeX
- Delivery markers are written
inline
[[pause]] [[breath]]
[[emphasis]]

The screenshot shows the RealMe Studio interface. At the top, it says "1 • Draft the narration". Below this, a note states: "Nothing is synthesized in this step, so it costs one cheap model call. The two context boxes matter more than the model you pick — they are what makes the script sound like you rather than like a textbook." There is a section for "Slide deck (PDF, PPTX or ODP)" with a "Choose File" button and "No file chosen" text. Below that is a "Course context" field. The interface has a header with "RealMe Studio" and three status indicators: "MEDIA TOOLS READY" (green), "NO GEMINI KEY" (orange), and "PLACEHOLDER VOICE ONLY" (yellow). There are three tabs: "Lecture" (selected), "Podcast & Debate", and "Twin Setup". The main content area shows a text box with the text: "Dry and precise. Worked examples over abstraction. I say 'notice that' a lot and I never read a bullet aloud." Below this is a "Script writer" dropdown menu set to "Gemini" and a "Model" dropdown menu set to "gemini-3-flash (free tier)". A blue "Draft narration" button is at the bottom of this section. Below the main content area is a section titled "2 • Read it. Fix it. Then render." with a note: "Every edit here is free; every word below this point costs money and minutes. Fix acronyms, cut padding, add the aside you always make in class." On the right side, there is a panel titled "Rendering lecture — done" showing a list of segments with timestamps and durations: 05:38:12 Preflight, 05:38:12 segment 1: 15.1s measured, 05:38:26 segment 2: 34.7s measured, 05:38:55 segment 3: 34.6s measured, 05:39:24 segment 4: 34.2s measured, 05:39:52 Assembling master, 05:39:52 Verifying, 05:39:53 PASS (placeholder content), and 05:39:53 Finished.

A WORKED EXAMPLE

One deck, start to finish


Study Designs in Epidemiology



1

INPUT a four-slide deck, test_slides.pdf

Study Designs in Epidemiology



1

OUTPUT test_slides.mp4 — 5m 29s, 1080p, with captions

The video ships beside this file. Notice the highlight in the frame: whichever point is being spoken is the one emphasised on screen, and the cut list that does it is what later lets a lecture be split at exact slide boundaries.

What else it does

Narrate a document

Notes, a paper, a handout — read aloud with captions. The fastest loop in the system: check a pronunciation, or hear how a paragraph lands, in seconds rather than hours.

Podcast & debate

You and a second voice discuss a paper or a topic. Short turns, one point each — long turns read as alternating lectures rather than a conversation.

Split & join

Cut a recorded lecture at exact slide boundaries, or build a new one from sections of several. Each part stays an editable lecture of its own.

Revising a deck you have already recorded is the same idea one level down: give the Lecture tab the new version and it compares slide by slide — slides that only moved keep their notes, changed ones are flagged, new ones get fresh notes. Only what actually changed is spoken again.

BE REALISTIC

What it costs

Three voices, one bill. Only the hosted one costs money, and it costs almost no time; the local two cost time and nothing else.

about \$1

of Gemini for a 50-minute lecture

~5 hours

of your own CPU, for the same lecture

seconds

to hear one line before committing

How not to waste it

- Draft with a stock voice first
It tells you whether the timing and the writing work.
- Fix one sentence, re-record one sentence
Renders are keyed on the exact text and voice.
- Let it run unattended
Closing the browser does not stop it.
- Build for a GPU if there is one
One command; the setup check says whether it is worth it.

GETTING STARTED

Where to begin

```
00_Windows\1_Install.bat    → environment and engine
00_Windows\2_Set_Key.bat    → paste your Gemini key
00_Windows\4_Start_Studio.bat → record your voice, then a first deck
```

The guide is START_HERE.md in the project folder, and the same material is in the app's Help tab, beside the thing it describes.

When something is wrong, three commands answer most of it:

```
realme setup — what is missing and how to fix it
realme key --where — which key is being read
realme doctor — which voices can actually run
```

A synthetic voice that sounds like a named person should be labelled as one. RealMe records a consent statement alongside your voice and writes a disclosure line into every project. Keep both.