RealMe Studio

Your slides, narrated in your own voice

Record yourself once. Any slide deck comes back as a finished lecture video — narration, captions and chapters — spoken in a voice cloned from that recording.

Guide v4 · October 2026

What it makes

One recording of your voice, four kinds of output

1 Lectures

A slide deck becomes a narrated video with captions and chapters. Revise the deck later and only the changed slides are spoken again.

2 Read aloud

Paste a paragraph or give it a document. It comes back as audio with captions — also the quickest pronunciation check.

3 Podcasts

A topic or a paper becomes a two-voice debate, dialogue or interview between you and a second speaker.

4 Tools

Play a video between two slides, split a long lecture into parts, or join sections from different lectures.

Every word is yours to read and edit before anything is spoken.

What you need

You need

Windows 10 or 11 and Miniconda or Anaconda (the one thing a zip cannot carry is Python)

A Gemini API key from Google AI Studio — it drafts the narration; for the Gemini clone, link the project to billing

Your slides as a PDF PowerPoint: File → Save As → PDF

A quiet room and a decent microphone only for a cloned voice: one minute of your speech

You do not need

A graphics card

A compiler

Audio or video software

More than about 2 GB of disk

Install in three double-clicks

All in the RealMe folder → 00_Windows

1 **1_Install.bat**

Installs RealMe, ffmpeg and three free draft voices — a few minutes. Then asks: your own cloned voice? Gemini, the free local clone, or none (draft voices).

2 **2_Set_Key.bat**

Paste your Gemini key. Stored in your own data folder (%USERPROFILE%\RealMeStudio), so updates never touch it.

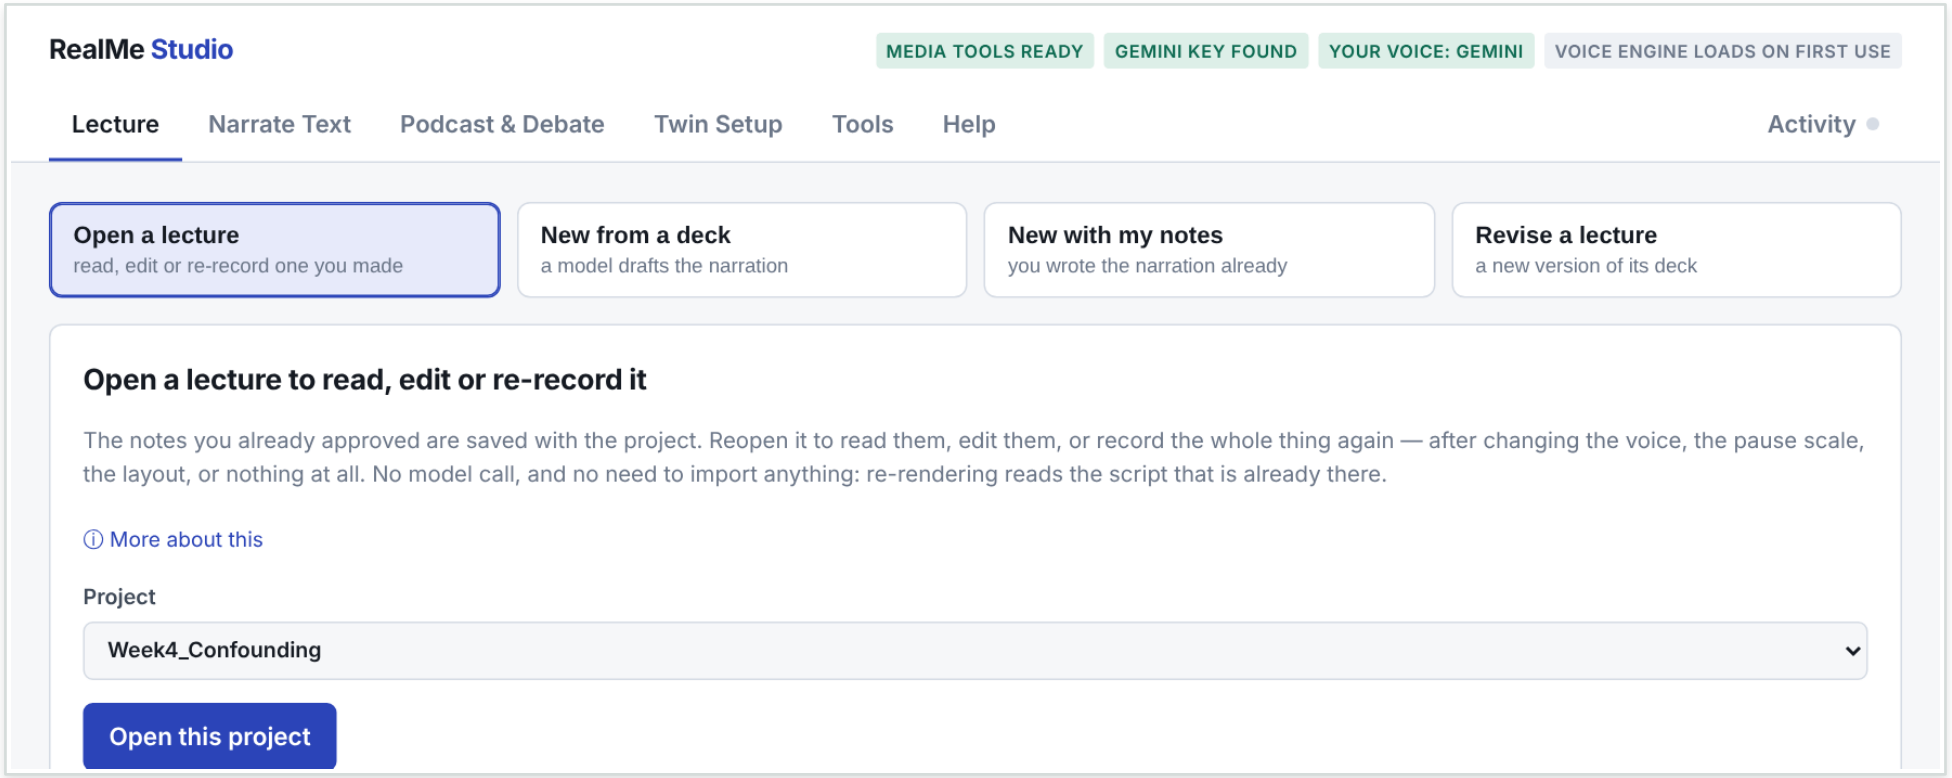
4 **4_Start_Studio.bat**

Opens the Studio in your browser at 127.0.0.1:8000. Nothing is online — the page is served by your own computer.

From then on, only step 4.

The free local clone, qwen3cpp: nothing to pay and nothing leaves the computer — it just takes longer to render. Ask Xinhua Yu, the RealMe developer, for a package that includes it, ready to run. Step 3: **3_Get_Voice_Engine.bat** will allow you to install and compile the realme-qwen3cpp engine.

The Studio at a glance



Six tabs

Lecture · Narrate Text · Podcast & Debate · Twin Setup · Tools · Help

Status pills

Green is ready. Anything missing says, in plain words, what to do next.

Activity

Says what is happening: drafting, recording slide 3 of 12, creating the video.

Record your voice once

Twin Setup → Your voice · only if you chose a cloned voice

30 to 60 seconds of you reading something you have the text of

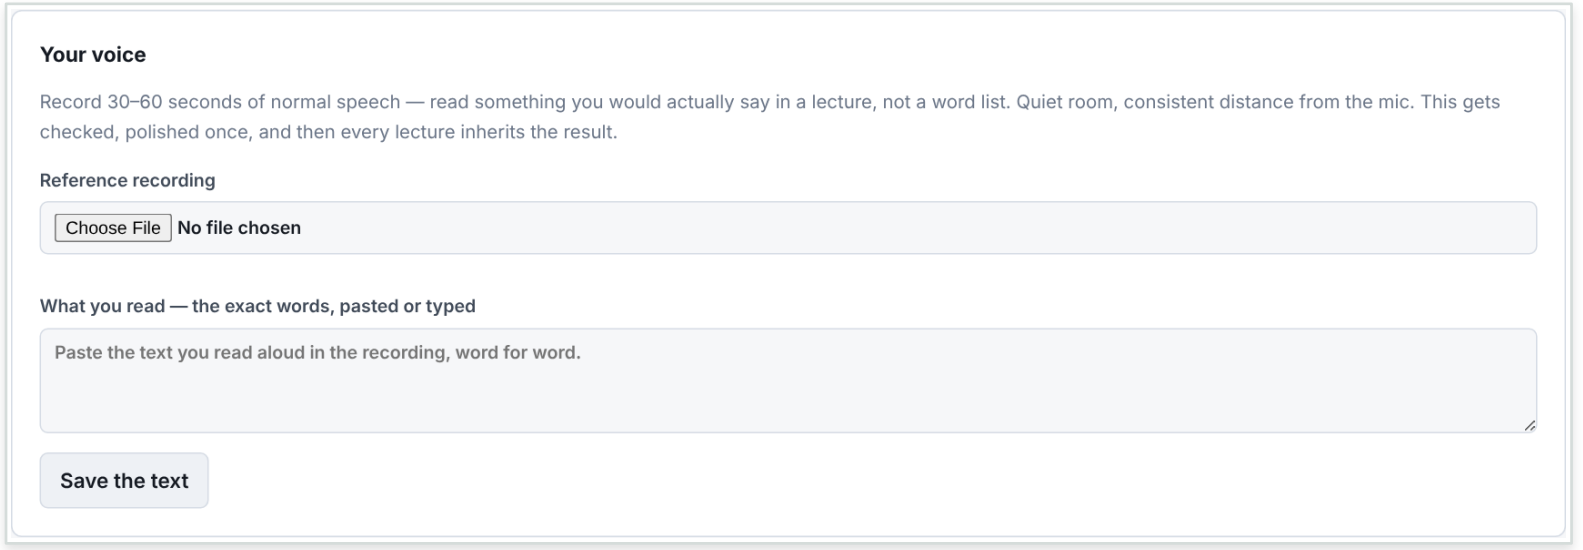
What you read paste the exact words, press Save the text

Read the way you lecture not slowly, not a performance: the clone copies what you give it

One sitting, one room, one microphone

The page checks it length, level and noise — if it says record again, record again

Save this voice when it is right



Create your cloned voice

Twin Setup → Gemini cloned voice · if you chose Gemini

1 **Confirm a paid project**

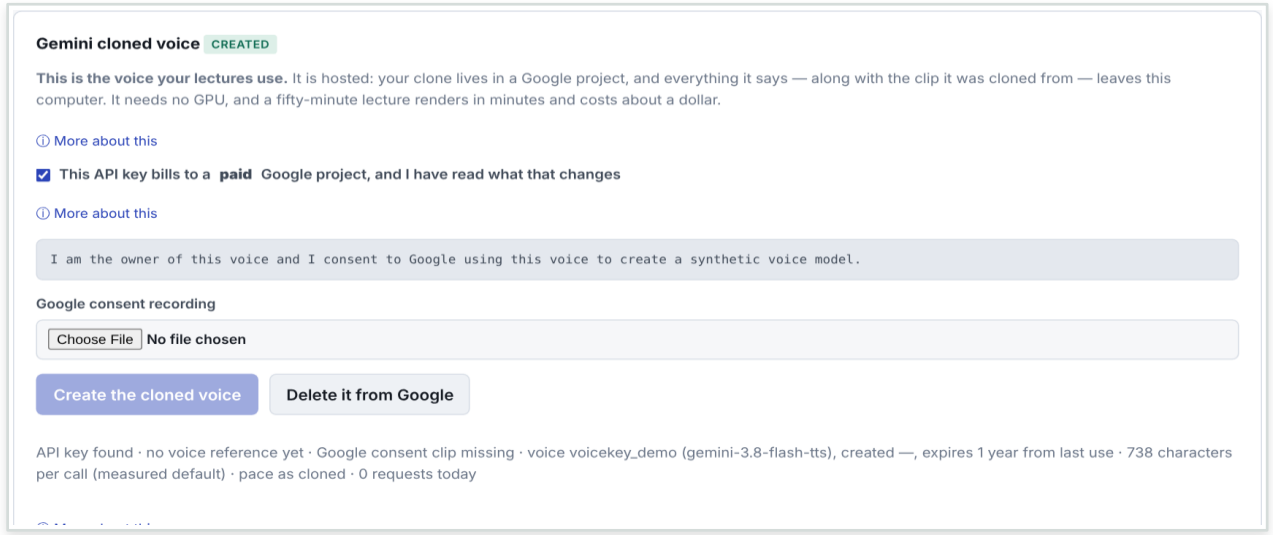
On the free tier Google may review what you send — and your voice is personal information.

2 **Record the consent sentence**

Google requires its exact wording. Same room, same microphone.

3 **Create the cloned voice**

Greyed out until everything is there — the line under it says what is missing. Then every lecture speaks in it.



Cheap first, expensive last

Every step catches a mistake before the costly one

Draft

one model call

seconds

›

Read & edit

free

where quality is decided

›

Hear a line

seconds

check pronunciation

›

Draft render

free voice

check timing

›

Final render

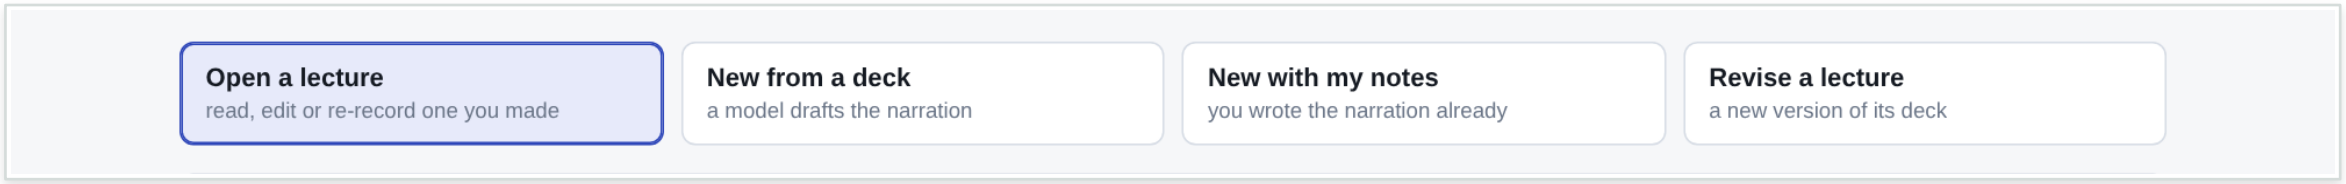
your voice

≈ $1 per hour of lecture

Iterate with the draft. Record once in your own voice.

Start a lecture

Lecture tab — four ways in



Open a lecture

read, edit or re-record one you made

New from a deck

a model drafts the narration

New with my notes

you wrote the narration already

Revise a lecture

a new version of its deck

Slides go in as a PDF. A PowerPoint file works only where LibreOffice is installed. Without it, RealMe says so at the upload and asks for the PDF version — PowerPoint: File → Save As → PDF.

Draft from a deck

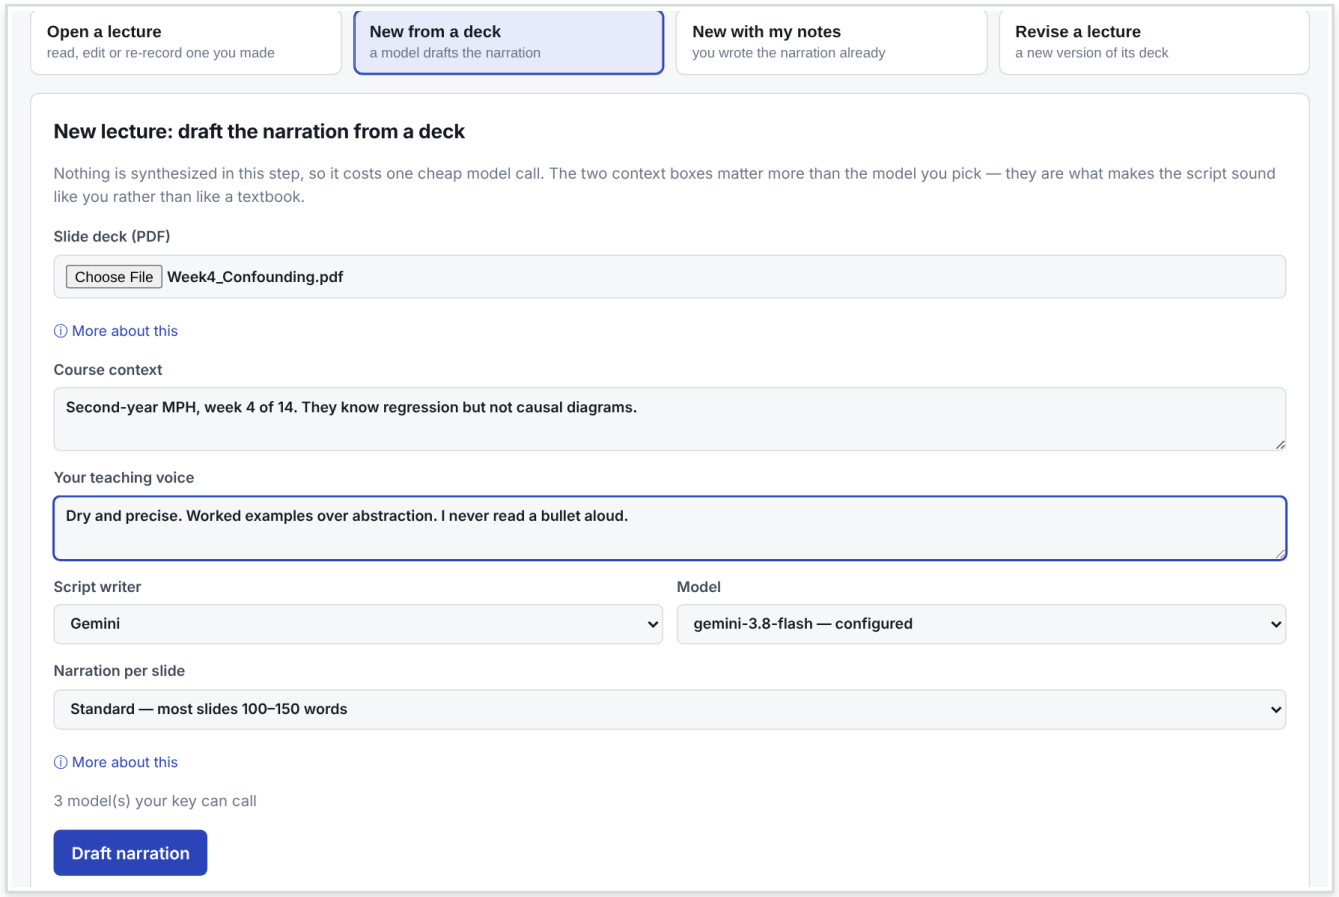
New from a deck → two boxes that matter most

Course context who the students are and where they are in the term

Your teaching voice how you talk — “dry and precise, worked examples, I never read a bullet aloud”

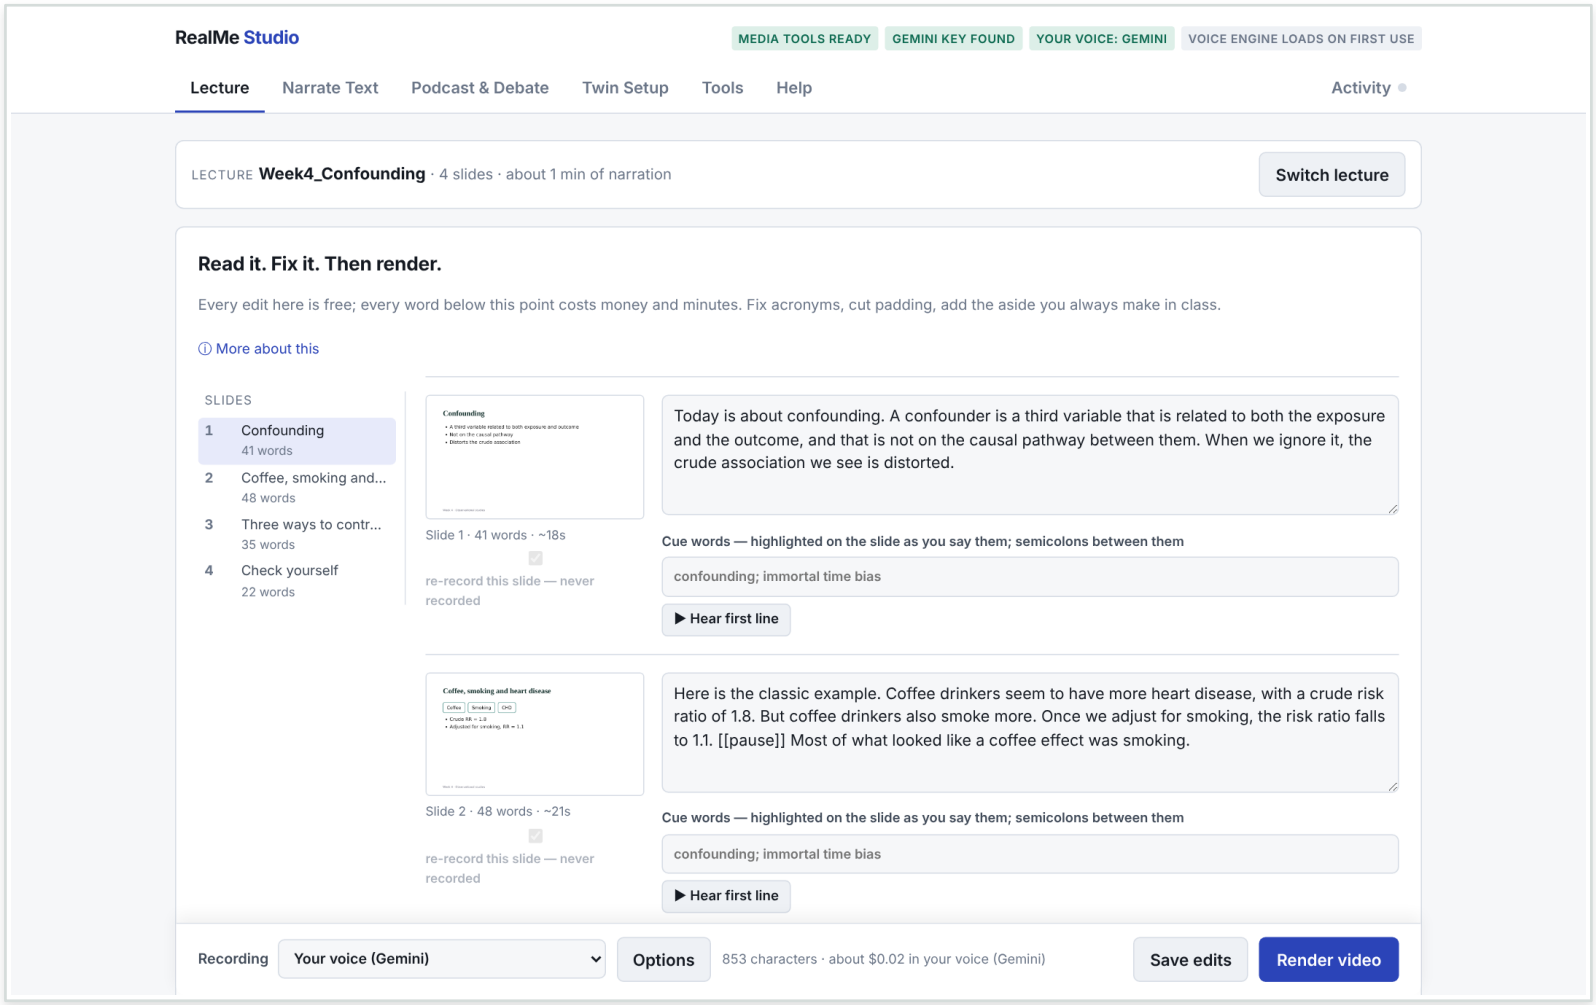
Narration per slide brief, standard or full

Draft narration a first version for every slide, in seconds



Read it. Fix it. Then render.

The editor — every edit here is free



Slide list jump to any slide

Narration with words and time per slide

Cue words highlighted on the slide as you say them; semicolons between

Hear first line a pronunciation check in seconds

Making it read well

Numbers

Write them as you would say them: “15.3 percent”.

Acronyms

Spelled letter by letter by default. A term said as a word gets one entry in Tools → Pronunciations, for every lecture after.

Equations

LaTeX is read as mathematics: \bar{x} becomes “x bar”, a/b becomes “a over b”.

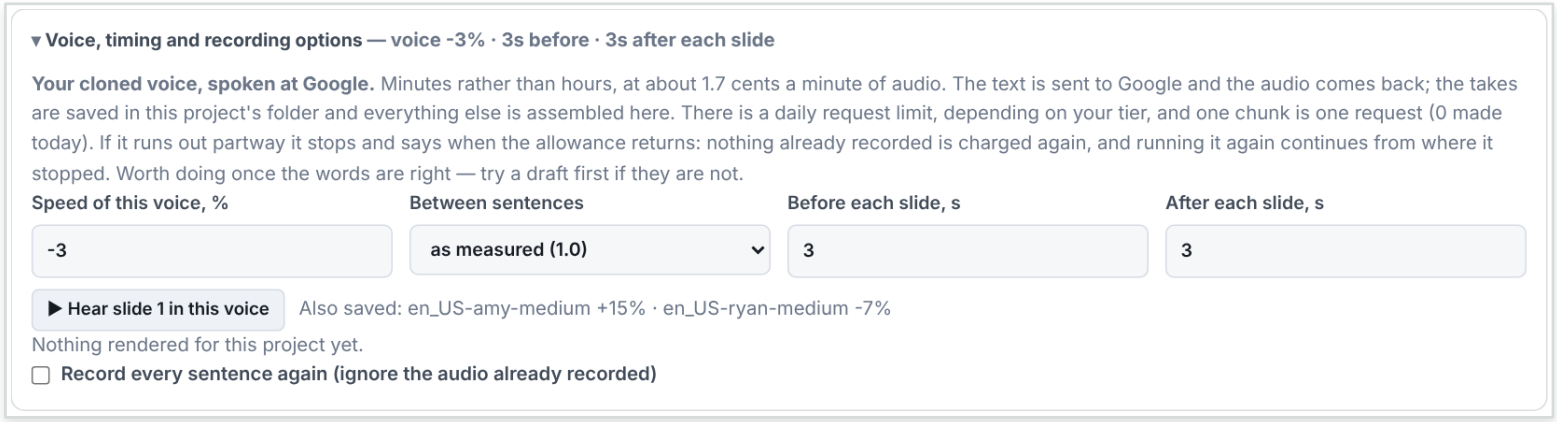
Delivery markers

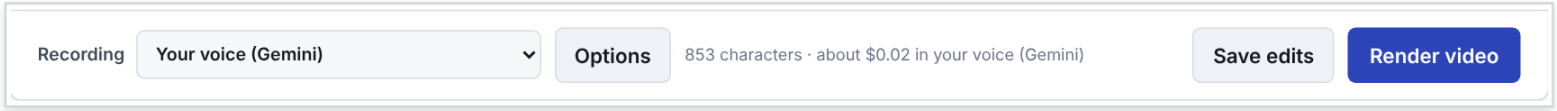
[[pause]] [[emphasis]] [[slow]] — instructions to the voice, never spoken, never in the captions.

Fix it in the text once; every later render inherits it.

Voice, timing and recording options

Below the slides — re-timing, never re-recording





Speed of this voice saved per voice, −30% to +30%

Between sentences tighter or more open pauses

Before / after each slide 3 s by default

Render bar Recording: your voice or a draft voice, in plain words; cost; Render video

Changing these is free: they re-time what is already recorded.

What a render gives you

`Week4.mp4`	the video
`Week4.srt` · `.vtt`	captions, two formats
`Week4_chapters.txt`	one chapter per slide
`cut list`	exact slide times, for splitting

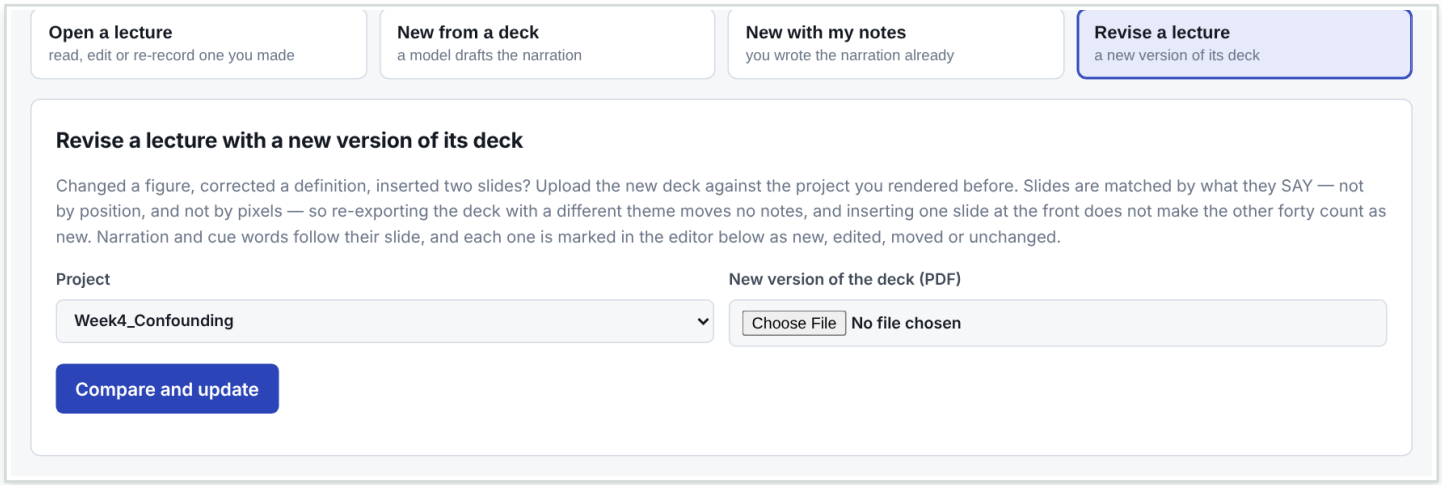
Draft first a draft-voice render is saved separately and can never overwrite the real one

Incremental sentences already spoken with the same words and voice are reused — fix one slide, pay for one slide

Stopped partway? start it again; it continues where it stopped

Revise a lecture

A new version of the deck, next term



Compared slide by slide by what the slides say, not their pixels

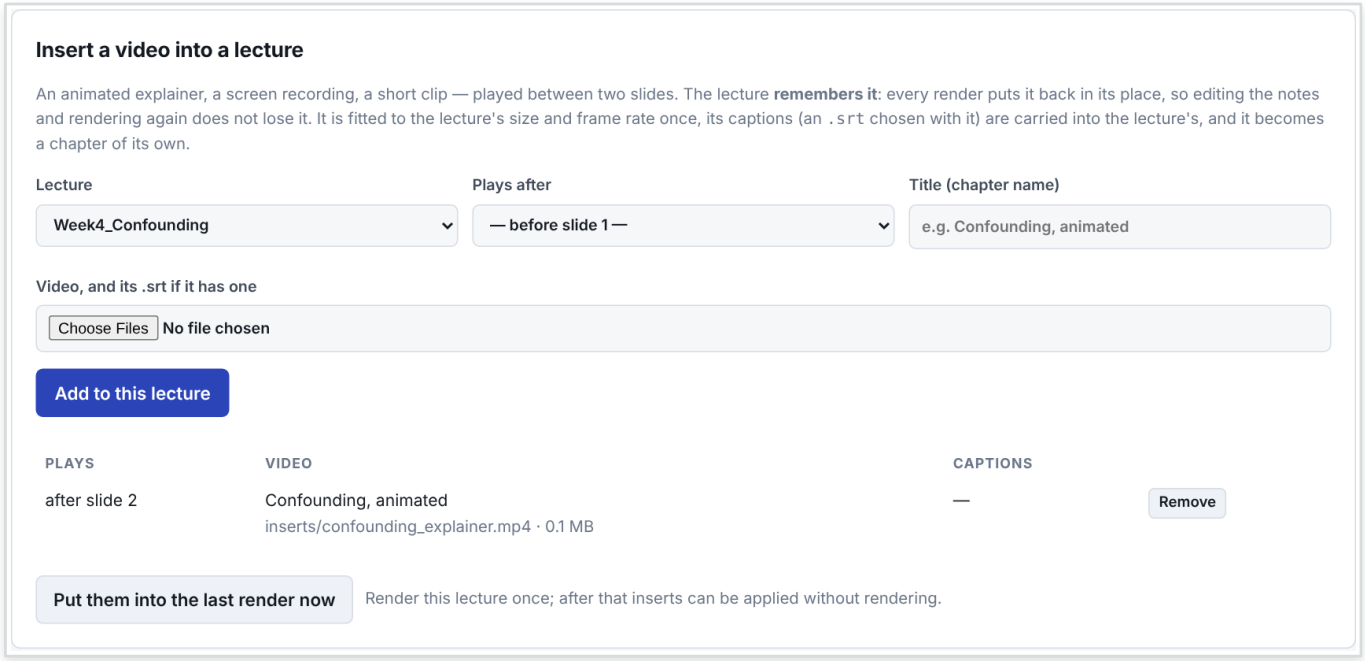
Narration follows its slide

Only the new is written and the edited checked

Only what changed is recorded usually a few cents

new edited moved unchanged

Tools: insert, split, join



Insert a video

An animated explainer or screen recording between two slides — remembered by the lecture; captions and chapters follow.

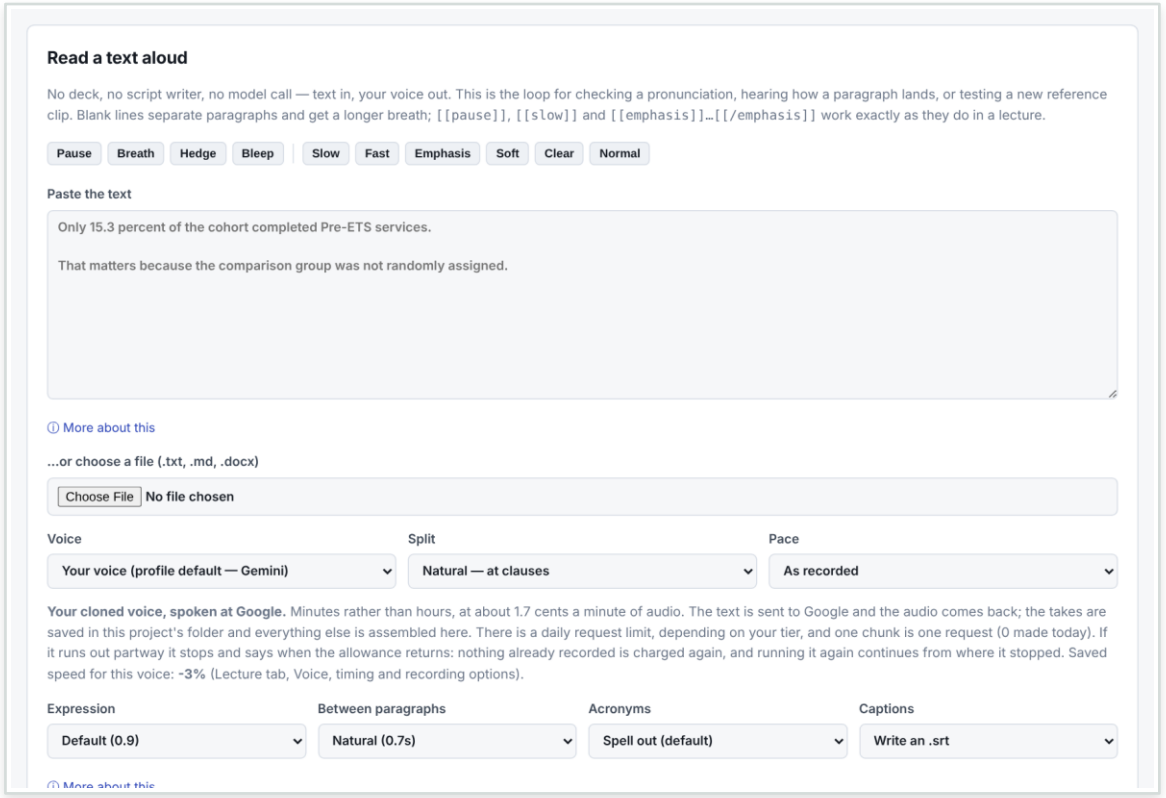
Split a video

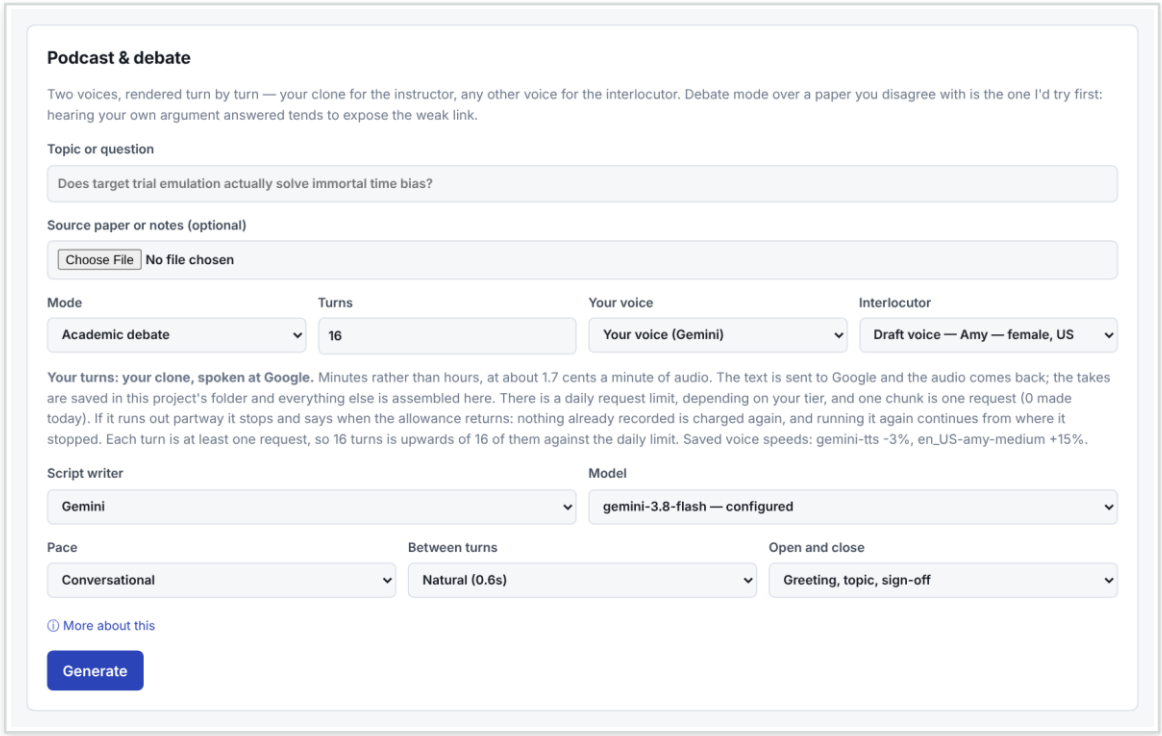
Cut a long lecture into parts at exact slide boundaries.

Join videos

Put sections from different lectures together.

Read aloud, and podcasts





Narrate Text paste text or choose a .txt, .md or .docx file — audio with captions

Podcast & Debate a topic or paper becomes a debate, Socratic dialogue or interview, in your voice and a second one

Cost, privacy and consent

Goes to Google

The slide text, once, to draft. The narration, to be spoken in your clone — about $1 per hour of lecture. Use a paid project.

Stays on your computer

Your recordings, projects and videos, in %USERPROFILE%\RealMeStudio. Updates never touch it.

Tell your students

Twin Setup holds a short spoken disclosure. Say it in the first ten seconds of every lecture.

Prefer to pay nothing? qwen3cpp, the free local clone (ask Xinhua Yu, the RealMe developer), records at no cost and keeps even the narration on your computer — only a little bit slowly. The draft voices cost nothing at all.

> *“This lecture was narrated with a synthetic voice model of the instructor, from a script the instructor wrote and approved.”*

Where to go from here

1	2	3	4	5	6
Install	Record your voice	Draft	Read & fix	Hear a line	Render

Help tab inside the Studio covers every screen.

START_HERE.md in the RealMe folder is the full guide.

%USERPROFILE%\RealMeStudio holds your material; updates never touch it.

Stuck? Ask Xinhua Yu, the RealMe developer.

Start with a short deck — five or six slides — and listen.